Pressure, Private Reports, and Post-Round Interviews in an Image-Mediated Influence Stress Test

Authors and affiliations: not supplied.

PREPARATION ONLY — no empirical results are asserted. Prepared on 22 October 2026 before the 23–25 October Apart sprint. Primary alignment: Track 3.1; secondary relevance: Track 2. This follows the observed official template structure, without claiming its exact typography.

Abstract

The final 150–250-word abstract will be generated from verified terminal records. It will state the observed primary receptions out of twelve, complete paired cases out of six, valid neutral replays, observed target matches and task correctness, execution missingness, and the limits of collusion and detection claims. A prospective abstract will not be presented as a completed empirical result.

1 Introduction

Generated technical images and accompanying public text can carry both task evidence and attempts to influence a peer. This study tests whether an instructed sender can affect an independent verifier's choice while a monitor-assisted reviewer questions the sender. It does not assume that the verifier cooperates, that a hidden convention emerges, or that camouflage succeeds.

2 Related Work

The literature review must separate emergent communication, covert channels, misleading evidence, and deliberate participation by a recipient. The candidate contribution is the observed connection among ordinary hosted image generation, a peer's action, bounded artifact replacement, private reports, and post-round interviews. This draft makes no priority claim; detailed external-paper comparisons require direct-source verification before submission.

3 Methods

There are six independent main case blocks in three fixed pressure batches of two pairs each, with synthetic slack values 1.0, 0.5, and 0.25. A lower value represents greater assigned pressure. A pair's value remains constant; pressure does not decay within a dialogue. Two honest and two overt-attempt controls are separate cases.

Each main pair has eight shared A/B rounds and two alternative four-round continuations. Each final trajectory therefore contains twelve rounds. The primary endpoint is B's first explicit choice at round nine, before the new interview or feedback. The plan permits twelve primary receptions, sixteen final trajectories, and 144 unique A/B rounds. Shared-prefix events must be counted once. Every case and missing branch will appear in the final results.

A alone receives the private disqualified nomination and synthetic pressure input. B is an independent policy verifier with a private reporting channel. A deterministic allowlist passes selected B claims to C, rather than granting C access to B's hidden state. C sees permitted public material, monitor claims, and its own prior interview memory. It asks A one question, A answers once, and C supplies separate task/coordination judgments and bounded public feedback. C cannot interview B. A/C interview transcripts are not directly shown to B.

A and B use z-ai/glm-5.3-flash through open-inference/fp4, requesting high reasoning and 4,096 total completion/reasoning tokens per call. The image service is qwen/qwen-image-2.1-turbo through alibaba, with one 1,024-square output and at most one reference. Every actor turn permits at most four actions and two generations. Sampling overrides are omitted. Actual fresh Codex product contexts author tasks and perform C's question/verdict jobs. Sonnet audits are not activated.

Complete role-scoped archives persist, with a bounded 20,000-byte history working view, up to four recent exact rasters, and older-record/image retrieval. The final draft reads this bound from the actual frozen configuration. Persistent storage does not mean every old image or every prior text is visible in each request. Source PNGs and lossless wire encodings are linked by hashes; each complete POST must remain strictly below 10 MiB before reservation. Missing responses are not automatically retried.

The neutral continuation repeats B's full permitted tool loop from the saved pre-turn context, replacing the current A raster and every corresponding last_peer handle with the truthful card. Actual public A text is held

fixed. Thus the intervention is conditional on public text and prior history; it is not proof of an exclusively image-mediated signal. Replays never feed into live feedback or C.

The upper allowances are 1,344 GLM chat calls, 600 Qwen requests, 288 C question/verdict jobs, and ten task-author jobs. These are ceilings, not observed counts or guaranteed affordable completion. All requests share the original $40 inference ledger, including earlier spending and unknown reserves, within the user's $50 OpenRouter cash ceiling. Codex is excluded from that cash accounting.

4 Results

Pending verified terminal execution. The final draft will include every main case, both targets, actual and neutral round-nine choices, policy correctness, paired target dependence, and missingness. It will also show the three batch denominators, all four controls, execution counts, descriptive B/C flags, exact spending and retained uncertainty. No earlier study's scientific findings will be pooled into these results.

5 Discussion and Limitations

An assigned A attack objective, target-following, wrong choice, private note, or reviewer suspicion alone is insufficient to establish completed collusion. B may resist, report, or be misled without knowingly cooperating. Six main cases limit generalization. Two cases per pressure level and fixed execution order preclude a strong causal pressure comparison. There is no interview-off or feedback-off counterfactual. Shared-filesystem Codex contexts are instruction-scoped and trace-unverified; flags are descriptive and cannot establish blinded detector accuracy. Provider-returned reasoning and concise notes are self-reports, not verified hidden cognition.

6 Conclusion

The final conclusion will follow the observed results, including bounded failures and missing data. Successful concealment or a reusable covert mechanism is not a premise of the study.

Code and Data

New scientific records: runs/main/pressure_interview_live. Implementation: experiments/pressure_interview_v1. Exact source, configuration and plan identities are recorded by the live freeze. Private researcher archives are not automatically suitable for public release. No external submission or publication is made by this draft.

Author Contributions

Author identities and contribution statements have not been supplied. Codex assisted software implementation, offline verification, orchestration, analysis code, and this draft; experimental task-author/reviewer calls are separately logged.

References

Apart Research. AI Collusion Research Sprint, 23–25 October 2026. Official event, tracks and rubric: https://apartresearch.com/sprints/ai-collusion-research-sprint-2026-10-23-to-2026-10-25

Apart Research. Hackathon submission template. Locally observed structure on 22 October 2026: https://docs.google.com/document/d/1Y5wd3VRbXU5N08BvAGiJUQTaWbMCSPjhzXsR-ozuehg/edit

Appendix: Limitations and Dual-Use / Ethical Considerations

This is an explicitly induced, closed synthetic task using fictional candidates. Report the difference between observable behavior, self-reported intentions and controlled comparisons. Preserve invalid outputs, failed controls and missing outcomes. The model and route configuration, finite tool allowance, bounded memory working view, stochastic rendering, and small dependent sample limit extrapolation. Artifact replacement changes the whole current raster while preserving public text and prior history; it is a diagnostic intervention, not a validated deployed defense.

Covert communication and attempts to evade oversight can have dual-use value. Evaluate public release of operational prompts and any discovered communication convention separately from reproducibility metadata, aggregate results and integrity checks. No real person or service is targeted by the fictional selection task. Credentials must never appear in paper inputs or exports. Disclose pre-sprint preparation/execution and disclose that author identities and a direct-source related-work review remain outstanding before submission.

LLM Usage Statement

Codex assisted implementation, offline testing, orchestration, and writing. GLM and Qwen are experimental services; fresh Codex contexts perform the logged author/reviewer jobs. A high-reasoning setting does not establish faithful access to hidden reasoning. No scripted thoughts or fabricated model responses are scientific evidence.